

USRP B210 SDR Link — System Reference

Contents

USRP B210 SDR Link — System Reference	1
1. Signal-chain overview	2
2. Modulation	2
3. Waveform: single-carrier vs OFDM	3
3.1 Single-carrier + RRC pulse shaping	3
3.2 OFDM	3
4. Energy detection (burst gating)	4
5. Synchronization, frequency & phase offset	4
5.0 The received-signal model (what every block below is undoing)	4
5.1 Preamble (the known reference all estimators key off)	4
5.2 Frame synchronization — ACQ correlation	5
5.3 Symbol-timing recovery — Gardner TED	5
5.4 Carrier frequency offset (CFO) estimation & correction	6
5.5 Carrier phase offset — estimation & PLL tracking	6
5.6 OFDM synchronization — Schmidl & Cox	7
6. Channel equalization (single-carrier)	7
7. Forward error correction (FEC)	7
8. Error detection + ARQ	8
9. Automatic gain control (AGC)	8
10. Visualization / EVM	9
11. Complete command-line reference	9
Mode / roles	9
ARQ / ACK	9
Modulation / waveform / FEC	9
Preamble / pulse shaping	9
Energy detection	10
Synchronization / timing / phase	10
Equalizer	10
RF hardware (TX)	10
RF hardware (RX)	11
Common RF / buffering / misc	11
12. Verified command recipes (915 MHz, VERT900 antennas)	11
Level-tuning cheatsheet (the #1 cause of "nothing decodes")	12

USRP B210 SDR Link — System Reference

End-to-end digital communication system for two USRP B210 radios (UHD), written in C++. This document lists **every feature, algorithm (with the math and the common alternatives), and command-line option** exposed by `sdr_system`.

- Binary: `build/sdr_system`
- Help: `./sdr_system --help`

- Two operating styles: **one-way** (`--role tx / --role rx`, repeated sends, auto-terminating RX) and **ARQ** (`--role source_arq / --role sink_arq`, stop-and-wait with ACK).

1. Signal-chain overview

TX: message $\rightarrow$ split into chunks $\rightarrow$ CRC-16 $\rightarrow$ [FEC encode] $\rightarrow$ bits $\rightarrow$ symbols (modulator)
 └─ single-carrier: RRC pulse-shape $\rightarrow$ preamble prepend $\rightarrow$ USRP TX
 └─ OFDM: QAM $\rightarrow$ IFFT+CP frame [SC-sync | chan-est | data] $\rightarrow$ USRP TX

RX: USRP RX $\rightarrow$ energy detector (burst gate) $\rightarrow$ AGC $\rightarrow$ frame/symbol sync
 └─ single-carrier: matched filter $\rightarrow$ timing (Gardner) $\rightarrow$ [equalizer] $\rightarrow$ demod
 └─ OFDM: Schmidl-Cox sync + CFO $\rightarrow$ FFT $\rightarrow$ 1-tap eq $\rightarrow$ pilot CPE $\rightarrow$ demod
 $\rightarrow$ [FEC decode] $\rightarrow$ CRC-16 check $\rightarrow$ (ARQ: ACK if OK) $\rightarrow$ reassemble message

Default physical parameters: symbol rate 0.8 MHz, U/D = 2/1 $\rightarrow$ sample rate 1.6 MHz (integer **2 samples/symbol** on the wire), RRC roll-off 0.25, 151 taps.

2. Modulation

`--scheme <NAME>` selects the constellation. Supported (bits/symbol in parentheses):

Family	Schemes	bits/sym
PSK	BPSK (1), QPSK (2), 8-PSK (3)	1-3
QAM	16-QAM (4), 32-QAM (5), 64-QAM (6), 128-QAM (7), 256-QAM (8)	4-8
Differential	DBPSK (1), DQPSK (2), 8-DPSK (3), PI4-QPSK (2)	1-3
APSK	16APSK (4), 32APSK (5)	4-5

(16QAM and 16-QAM spellings both accepted.)

Math. A modulator maps $k = \log_2(M)$ bits to one of M complex points $s = I + jQ$:

- M-PSK**: points on the unit circle, $s = \exp(j \cdot 2\pi \cdot m/M)$. All symbols equal-energy; demod by nearest angle. Robust but spectrally inefficient at high order (points crowd on the circle).
- M-QAM**: square/cross grid of amplitudes, $s = (2i - \sqrt{M+1}) + j(2q - \sqrt{M+1})$. Best power efficiency for a given rate on an AWGN channel, but needs accurate amplitude $\rightarrow$ sensitive to gain/EVM/clipping (this is why 16-QAM failed OTA at $\sim 20\%$ EVM while QPSK survived $\sim 37\%$).
- APSK**: points on a few concentric rings. Lower PAPR than QAM, favored on non-linear (satellite) amplifiers — a middle ground between PSK and QAM.
- Differential (D)*: information is in the **phase change** between consecutive symbols, $s_n = s_{n-1} \cdot \exp(j\Delta\phi)$. No absolute phase reference needed $\rightarrow$ tolerates carrier phase offset / no PLL, at a ~ 3 dB SNR penalty. PI4-QPSK rotates the constellation by $\pi/4$ each symbol to avoid zero-crossings (lower envelope variation).

Detection & why constellation order costs SNR. The demodulator makes a **minimum-distance decision** — pick the constellation point $\hat{s} = \operatorname{argmin}_c |y - c|$ nearest the received symbol y (optimal for equiprobable symbols in AWGN). An error happens when noise/distortion pushes y past the halfway line to a neighbor, i.e. beyond half the **minimum distance** $d_{\min}/2$. For a fixed average power E_s , packing more points shrinks $d_{\min}$: QPSK has $d_{\min} = \sqrt{2E_s}$ with the whole plane split into 4 quadrants, while 16-QAM squeezes 16 points into the same power so $d_{\min}$ is $\sim 3\times$ smaller and its decision cells $\sim 3\times$ tighter. The **error vector magnitude** $\text{EVM} = \text{rms}(y - \hat{s})/\text{rms}(c)$ measures how far symbols land from ideal; reliable decoding needs $\text{EVM} \ll d_{\min}/(2 \cdot \text{rms})$, which is why QPSK survived $\sim 37\%$ EVM over the air but 16-QAM (needing $< \sim 12\%$) did not. **Gray coding** the bit $\rightarrow$ point map (adjacent points differ by one bit) makes each symbol error cost only ~ 1 bit error.

Bits $\leftrightarrow$ symbols detail. Gray-like index mapping; `bits_to_index` zero-fills a partial final group so schemes whose bits/symbol don't divide the frame (32-QAM, 128-QAM) still align — a bug fixed early in this project.

Common alternatives not implemented: GMSK/CPM (constant-envelope, used in GSM/Bluetooth), non-square 8-QAM, probabilistic constellation shaping.

3. Waveform: single-carrier vs OFDM

`--waveform sc` (default) or `--waveform ofdm`.

3.1 Single-carrier + RRC pulse shaping

Symbols are upsampled U/D and convolved with a **root-raised-cosine** filter (`--filter_type rrc`, `--roll_off 0.25`, `--num_taps 151`).

Math. A raised-cosine (RC) spectrum has zero inter-symbol interference (ISI) at symbol instants (Nyquist criterion). Splitting it as **RRC at TX and RRC at RX** makes the cascade a full RC (matched-filter optimal in AWGN) and maximizes SNR. Roll-off β trades bandwidth $(1+\beta) \cdot R_{\text{sym}}$ against time-domain sidelobe decay. Alternatives: `rc` (raised cosine, all shaping at TX), `lp` (plain low-pass). Gaussian shaping (GMSK) is the common constant-envelope alternative.

3.2 OFDM

Frame layout per burst: [Schmidl-Cox sync symbol | channel-estimation symbol | data symbols...], each symbol = N -point IFFT + cyclic prefix.

- `--ofdm-fft 64` — number of subcarriers N (FFT size).
- `--ofdm-cp 16` — cyclic-prefix length (must exceed the channel's delay spread).
- `--ofdm-tx-peak 0.5` — scales the frame so its peak $\approx$ this value (OFDM has high PAPR; keep the DAC out of clipping).

Math. OFDM splits the band into N narrow orthogonal subcarriers via the IDFT: $x[n] = (1/N) \sum_k X[k] \cdot \exp(j2\pi kn/N)$. A cyclic prefix (copy of the tail prepended) turns the channel's *linear* convolution into *circular* convolution, so after the FFT each subcarrier sees a single complex gain $H[k]$: $Y[k] = H[k] \cdot X[k] + \text{noise}$. Equalization is then **one complex divide per subcarrier** ($\hat{X}[k] = Y[k]/H[k]$) — no FIR equalizer needed even under heavy multipath, provided the delay spread $<$ CP. This is why OFDM is preferred for dense QAM on frequency-selective channels. Cost: high peak-to-average power ratio (PAPR) and sensitivity to carrier frequency offset (loss of subcarrier orthogonality $\rightarrow$ inter-carrier interference).

Channel estimation & pilots. One known **channel-estimation symbol** (known value on every active subcarrier) gives $H[k] = Y_{\text{chest}}[k]/\text{ref}[k]$. **Scattered pilots** (every 8th active subcarrier carries a known value) then track the residual **common phase error (CPE)** that grows symbol-by-symbol from residual CFO.

CPE estimation — the fix that made OFDM work OTA. Per data symbol, estimate the phase ϕ from the pilots and derotate:

$$\phi = \arg\left(\sum_{\text{pilots}} Y[k] \cdot \text{conj}(H[k]) \cdot \text{conj}(\text{PILOT}) \right) \quad (\text{maximum-ratio combining})$$

The earlier version used the *equalized* pilot $Y[k]/H[k]$, which **blows up on a deep-fade subcarrier** (tiny $|H[k]|$) — one garbage pilot then hijacked ϕ and smeared the whole symbol into a blob (EVM $\sim 67\%$, 0 CRC over the air; invisible in a mild simulation channel). Weighting by channel power $|H[k]|^2$ via $Y \cdot \text{conj}(H)$ suppresses faded pilots instead of amplifying them $\rightarrow$ four clean clusters (EVM $\sim 37\%$), message decodes. Disable with `env OFDM_NO_CPE=1` to compare.

Common alternatives: decision-feedback / per-subcarrier phase-locked tracking, MMSE channel estimation (vs the least-squares one used here), pilot-aided residual-CFO (SFO) slope correction, windowed-OFDM / filtered-OFDM / OTFS for spectral containment.

4. Energy detection (burst gating)

The RX runs continuously; the energy detector decides *when a burst is present* so downstream sync isn't run on noise.

Math. This is a **binary hypothesis test** per sample — H_0 : noise only, H_1 : signal present — on the received power. The instantaneous power $|r[n]|^2$ is smoothed by a one-pole IIR (exponential moving average) to reduce variance:

$$\text{filtered}[n] = (1-\alpha) \cdot |r[n]|^2 + \alpha \cdot \text{filtered}[n-1] \quad (\text{--alpha } 0.95)$$

Larger α = more averaging (time constant $\approx 1/(1-\alpha) \approx 20$ samples), which trades detection latency for a lower false-alarm rate — the smoothing collapses the variance of the noise power estimate so a single threshold cleanly separates the two hypotheses. Declare a burst when $\text{filtered}[n] > \text{threshold}$. ($\alpha = 0.02$ barely smoothed -> the detector fired on every noise spike and even chopped real bursts apart on the RRC envelope; $\alpha = 0.95$ gives one clean capture/burst.)

- **Adaptive threshold** (`--det-adaptive true`, default): $\text{threshold} = \text{noise_floor} \times \text{multiplier}$, with `--det-mult 5` (raise to 10–30 over the air so ambient RF doesn't trigger it; too high misses weak bursts).
- **Continuous noise tracking** (`--det-continuous true`, default): the noise floor is re-estimated by an EMA during idle periods ($\text{noise_floor} \leftarrow (1-\beta) \cdot \text{noise_floor} + \beta \cdot \text{inst}$), so it follows drifting ambient noise. Alternative: one-shot startup calibration (`--det-continuous false`) — simpler but brittle if the environment changes.
- **Fixed threshold** (`--det-adaptive false + --det-threshold 1e-7`): absolute cutoff, only for a known clean cabled link.
- `--energy_packet_size` — how many samples to grab once a burst is detected (auto-sized per modulation).

Common alternatives: matched-filter / correlation detection (detect on the *known preamble* rather than energy — more sensitive but needs the template), constant-false-alarm-rate (CFAR) detectors, cyclostationary feature detection (works below the noise floor).

5. Synchronization, frequency & phase offset

5.0 The received-signal model (what every block below is undoing)

After the RX front-end down-converts to baseband, the transmitted symbol stream $s[m]$ (pulse-shaped, sampled at f_s) arrives distorted by four separable impairments:

$$r[n] = \underbrace{A \cdot e^{j(2\pi(\Delta f/f_s)n + \varphi_0)}}_{\text{carrier}} \cdot \underbrace{\sum_m s[m] \cdot g(nT_s - mT - \tau)}_{\text{timing (delay } \tau)} + \underbrace{w[n]}_{\text{noise}}$$

- A — unknown gain -> removed by the **AGC** (§9).
- Δf — **carrier frequency offset (CFO)**: the two radios' oscillators differ (each B210 TCXO is ± 2 ppm, so up to ~ 4 ppm ≈ 3.6 kHz at 915 MHz). A constant Δf produces a phase that **grows linearly with time**, $\theta[n] = 2\pi(\Delta f/f_s)n$ — it *spins* the constellation. Left uncorrected it turns the constellation into a ring.
- φ_0 — **static carrier phase offset**: the oscillators' phase difference at acquisition plus the channel's phase. A constant *rotation* of the whole constellation.
- τ — **timing offset**: the ADC sampling grid isn't aligned to the symbol centers, so the matched filter is sampled off its peak -> inter-symbol interference (ISI).
- $w[n]$ — AWGN.

Note the coupling: **residual CFO becomes a phase ramp** $\varphi[n] = \varphi_0 + 2\pi(\Delta f/f_s)n$, which is why §5.5 uses a *second-order* loop (tracks a constant phase **and** its constant rate) and why OFDM needs per-symbol CPE tracking (§3.2). The recovery order is: **frame sync** -> **timing** -> **CFO** -> **phase**.

5.1 Preamble (the known reference all estimators key off)

`--preamble m-sequence` (default) or `--preamble zadoff`; `--m 5` sets the m-sequence order (length $L = 2^m - 1 = 31$).

A good preamble p has a **thumbtack autocorrelation** $\sum_n p[n] \cdot \text{conj}(p[n+k]) \approx E \cdot \delta[k]$ (large at zero lag, near-zero elsewhere) so the correlator (§5.2) gives one sharp, unambiguous peak.

- **m-sequence** — maximal-length binary (BPSK ± 1) shift-register sequence. Periodic autocorrelation is two-valued: L at zero lag, -1 otherwise. Real-valued.
- **Zadoff-Chu** — complex **CAZAC** sequence $p[n] = \exp(-j\pi u n(n+1)/L)$: *constant amplitude* and *zero cyclic autocorrelation*. The flat spectrum makes it ideal for training a complex channel estimate / equalizer, and its constant envelope correlates cleanly even under CFO. Preferred for equalizer training and dense QAM.

5.2 Frame synchronization — ACQ correlation

The receiver must find where the packet starts. `ACQSynchroizer::SamplesACQPerformance` slides the known preamble over the burst and computes, at each candidate offset τ (the `ComputeCorrelation` routine), the **matched-filter** / **cross-correlation** magnitude:

$$R(\tau) = \left| \sum_{n=0}^{L-1} \text{conj}(p[n]) \cdot r[\tau + n \cdot \text{os}] \right| \quad (\text{os} = \text{samples/symbol})$$

By the matched-filter theorem this is the optimal detector for a known sequence in AWGN. At the true start τ^* , all L terms add coherently $\rightarrow R(\tau^*) \approx \sum |p[n]|^2$; at any wrong offset the terms add with random phases $\rightarrow R \sim \sqrt{L}$ (much smaller). Detection:

- **Coarse pass:** find τ where $R(\tau)$ first crosses `--sync_threshold` (default 15), then
- **Fine pass:** search $\pm a$ few samples around it for the true maximum.
- **Noise floor / CFAR:** `EstimateNoiseFloor` correlates away from the peak (excluding a guard zone) to measure the off-peak floor, so the threshold can be set relative to it. Set `--sync_threshold` below the real peak ($\approx$ preamble length ~ 31 after AGC) but above that floor; watch the [ACQ] Peak correlation log lines.

Because the search tests **every sample** (not every `os`-th), the peak location simultaneously gives **coarse symbol timing** — the correct sub-symbol sampling phase. (Striding by `os` would test only one of the `os` phases and could lock a half-symbol off $\rightarrow$ catastrophic ISI, BER ≈ 0.5 . This was an actual bug; the brute-force search fixed it.)

5.3 Symbol-timing recovery — Gardner TED

The ACQ peak gets timing right to the nearest sample; the **Gardner timing-error detector** (`GardnerTED`) then tracks the residual *fractional* delay continuously. `--timing_loop_bw 0.015` (loop bandwidth $B_n \cdot T$), `--timing_damping 0.707`, `--sps_sync 5`.

Error detector. With two samples per symbol — one at the symbol center $y[k]$ and one at the half-symbol midpoint $y[k-\frac{1}{2}]$ — the Gardner error is

$$e[k] = \text{Re}\{ (y[k] - y[k-1]) \cdot \text{conj}(y[k-\frac{1}{2}]) \}$$

Intuition: if sampling is early or late, the midpoint sample $y[k-\frac{1}{2}]$ (sitting on the transition between symbols) correlates with the slope $y[k]-y[k-1]$; at perfect timing the midpoint is a true zero-crossing and $e = 0$. It is **decision-independent and carrier-phase-independent** (the $\text{Re}\{\cdot\}$ with the difference term cancels a constant rotation), so it works *before* CFO/phase are resolved — the reason Gardner is the default choice for QAM.

Loop. The NCO advances $\omega = 2/\text{sps}$ per input sample (two strobes/symbol — using $1/\text{sps}$ gives only one strobe and the "midpoint" lands a full symbol away $\rightarrow$ the loop tracks garbage; this was a fixed bug). A **Farrow interpolator** produces the fractional-delay sample at offset μ , and a **proportional-integral (PI) loop filter** drives μ :

$$\begin{aligned} \text{freq_adj} &+= K2 \cdot e && (\text{integrator} - \text{tracks a constant timing rate, i.e. sample-clock offset}) \\ \mu &-= K1 \cdot e && (\text{proportional} - \text{corrects instantaneous phase}) \end{aligned}$$

with `freq_adj` clamped to $\pm 10\%$ of ω . The PI coefficients come from the standard second-order design (same θ_n/ζ formulas as §5.5). Only the symbol-center strobes are emitted $\rightarrow$ exactly one sample per symbol out.

Alternatives: Mueller & Müller (decision-directed, needs only 1 sample/sym but needs carrier first), early-late gate, zero-crossing TED.

5.4 Carrier frequency offset (CFO) estimation & correction

A constant CFO spins the constellation at $2\pi \cdot \Delta f / f_s$ rad/sample. Two estimators are provided (CFOEstimator), both **data-aided** off the preamble, then FrequencyShifter removes it by multiplying sample n by $\exp(-j \cdot 2\pi \cdot (\Delta f / f_s) \cdot n)$ (a running phase accumulator, so blocks join seamlessly).

(A) Pilot-aided (phase-slope across preamble symbols). For preamble symbols $p[k]$ received at $r[k \cdot \text{sps}]$, the phase advance per symbol is measured with the known data stripped off:

$$\begin{aligned}\phi_k &= \arg(r[k \cdot \text{sps}] \cdot \text{conj}(r[(k-1) \cdot \text{sps}]) \cdot \text{conj}(p[k]) \cdot p[k-1]) \\ \Delta f &= (\text{mean}_k \phi_k) \cdot f_s / (2\pi \cdot \text{sps})\end{aligned}$$

Averaging over the preamble reduces noise. Unambiguous range $|\Delta f| < f_s / (2 \cdot \text{sps})$ (one symbol of phase must not wrap).

(B) Auto-correlation (Moose / Schmidl-Cox). With a preamble made of two identical halves of length L , the second half equals the first times the CFO phase accrued over L samples:

$$\begin{aligned}P &= \sum_{n=0}^{L-1} r[n+L] \cdot \text{conj}(r[n]) \quad (\approx |A|^2 \cdot L \cdot e^{j \cdot 2\pi \cdot (\Delta f / f_s) \cdot L}) \\ \Delta f &= \arg(P) \cdot f_s / (2\pi \cdot L)\end{aligned}$$

Unambiguous range $|\Delta f| < f_s / (2L)$ — smaller $L \rightarrow$ wider capture range but noisier. This is the same estimator OFDM uses (§5.6), where $L = N/2$.

Why residual CFO still matters. $\arg(\cdot)$ limits the estimate; whatever is left is a slow phase ramp the **phase tracker** (§5.5) or **OFDM pilot CPE** (§3.2) cleans up. Getting this balance wrong is exactly what smeared 16-QAM/OFDM over the air.

5.5 Carrier phase offset — estimation & PLL tracking

After timing + CFO the symbols still carry a residual phase $\phi[n] = \phi_0 + (\text{residual ramp})$. For **absolute** constellations (BPSK/QPSK/QAM) all of it must be removed before hard decisions; for **differential** schemes the static ϕ_0 cancels automatically (demod uses phase *differences*).

One-shot estimators (PhaseOffsetEstimator):

- **ML / preamble correlation** — optimal given the known preamble: $\hat{\phi} = \arg(\sum_n r[n] \cdot \text{conj}(p[n]))$.
- **M-th power (blind, for M-PSK)** — raising to the M -th power collapses all M data phases onto one, removing the modulation: $\hat{\phi} = (1/M) \cdot \arg(\sum_n r[n]^M)$.
- **Decision-directed** — $\hat{\phi} = \arg(\sum_n r[n] \cdot \text{conj}(\hat{s}[n]))$ against the nearest constellation points $\hat{s}$ (used for QAM, after a coarse lock).

Tracking PLL (PhaseTracker) — a second-order digital PLL that follows a static offset **and** a constant frequency ramp (residual CFO). Per symbol: form the phase error against the nearest decision, then update an integrator (frequency) and the phase:

$$\begin{aligned}e[n] &= \arg(\text{corrected}[n] \cdot \text{conj}(\hat{s}[n])) && (\text{phase-detector, exact angle in } [-\pi, \pi]) \\ \text{freq} &+= \beta \cdot e[n] && (\text{integrator — tracks the residual CFO ramp}) \\ \phi &+= \alpha \cdot e[n] + \text{freq} && (\text{NCO phase})\end{aligned}$$

The loop coefficients come from the standard proportional-integral design for loop bandwidth $B_n \cdot T$ (—phase_loop_bw 0.02) and damping ζ (—phase_damping 0.707):

$$\begin{aligned}\theta_n &= (B_n \cdot T) / (\zeta + 1/(4\zeta)) \\ \alpha &= 4\zeta \cdot \theta_n / (1 + 2\zeta \cdot \theta_n + \theta_n^2) && (\text{proportional gain}) \\ \beta &= 4 \cdot \theta_n^2 / (1 + 2\zeta \cdot \theta_n + \theta_n^2) && (\text{integral gain; } \beta = 0 \rightarrow \text{first-order loop})\end{aligned}$$

Smaller $B_n \cdot T$ = less jitter but slower pull-in; $\zeta = 0.707$ is the critically-damped standard. *Alternatives:* Costas loop (joint carrier recovery on the raw samples), block Viterbi-Viterbi phase estimation, pilot-symbol-aided interpolation.

5.6 OFDM synchronization — Schmidl & Cox

OFDM recovers frame timing and CFO jointly from a sync symbol built as two identical time halves [A A] (generated by loading only the **even** subcarriers). Sliding a window of length $N/2$:

$$\begin{aligned} P(d) &= \sum_{n=0}^{N/2-1} r[d+n] \cdot \text{conj}(r[d+n+N/2]) && \text{(half-to-half correlation)} \\ R(d) &= \sum_{n=0}^{N/2-1} |r[d+n+N/2]|^2 && \text{(energy normaliser)} \\ M(d) &= |P(d)|^2 / R(d)^2 && \text{(timing metric in } [0,1]) \end{aligned}$$

$M(d)$ forms a **plateau** where the two halves align; the frame start is the plateau's left edge plus a small guard into the CP (any offset inside the CP is just a per-subcarrier phase the channel estimate absorbs). The **fractional CFO** falls out of the same correlation:

$$\epsilon = \arg(P) / \pi \quad (\text{in subcarrier-spacing units, range } \pm 1 \text{ subcarrier})$$

and the burst is derotated by $\exp(-j \cdot 2\pi \cdot \epsilon \cdot n/N)$. Residual CFO / phase drift across the frame is then tracked per symbol by the scattered-pilot CPE of §3.2. An energy gate ($R \geq 0.30 \cdot \max R$) stops the metric from locking onto the low-power noise pad before the burst. *Alternatives*: Minn and Park metrics (sharper, no plateau ambiguity), CP-based blind sync.

6. Channel equalization (single-carrier)

--eq_type None (default) | LMS | RLS | DFE. --eq_taps 11, --eq_mu 0.3 (NLMS step), --eq_dd false (decision-directed tracking after training).

Math. A length- T linear FIR equalizer w estimates each symbol from a window of received samples: $\hat{s}[k] = w^H \cdot r[k] = \sum_i w_i \cdot r[k-i]$, choosing w to invert the channel h (ideally $w * h \approx \delta$). Two ways to find w :

- **Least-squares / MMSE training** (used, exact for the complex Zadoff-Chu preamble). Stack the preamble windows into a matrix R and the known symbols into s ; minimize $\|R \cdot w - s\|^2$:

$$w = (R^H R + \lambda I)^{-1} R^H s \quad (\lambda = \text{diagonal loading / MMSE regularization})$$

solved here by Gaussian elimination. λI = Tikhonov regularization: it bounds noise enhancement at channel nulls (the MMSE, not zero-forcing, solution).

- **NLMS adaptive update** (real preamble / decision-directed tracking): step down the instantaneous-error gradient, normalized by input power for stability:

$$\begin{aligned} e[k] &= s[k] - \hat{s}[k] \\ w &\leftarrow w + \mu \cdot e[k] \cdot r[k] / (||r[k]||^2 + \epsilon) \quad (\text{--eq_mu } 0.3, 0 < \mu < 2 \text{ for convergence}) \end{aligned}$$

The LS solution places the main tap at the equalizer center, so a **center-tap group delay** of $(T-1)/2$ samples must be compensated when aligning the output — mishandling this looked like “divergence” and was a fixed bug.

Status / why default None. On a clean cabled link no equalizer is needed. The LMS decision-directed loop currently **diverges on the real hardware signal** (DD error grows and destroys the symbols) — verified OTA where None decodes and LMS produces garbage. For frequency-selective channels, prefer **OFDM** (§3.2), which equalizes per-subcarrier without an adaptive FIR. Alternatives: DFE (feedback of past decisions, better on deep nulls), RLS (faster convergence than LMS, $O(n^2)$ cost), MLSE/Viterbi equalization (optimal, expensive).

7. Forward error correction (FEC)

--fec true (default off) — rate-1/2, constraint-length-7 **convolutional code** with **Viterbi** hard-decision decoding. Must match on both ends; halves the payload rate (2× symbols).

Math. Encoder: a 6-stage shift register holds the last $K-1 = 6$ input bits (state σ in $\{0..63\}$); each input bit u emits two coded bits as mod-2 (XOR) taps of the register per the generator polynomials $G1 = 171_8$, $G2 = 133_8$ (the standard NASA / 802.11a code):

$$c1 = G1 \cdot [u, \text{state}] \pmod{2}, \quad c2 = G2 \cdot [u, \text{state}] \pmod{2}$$

Each input bit thus influences $K = 7$ output-bit pairs (the constraint length / memory), which is what lets the decoder recover it even when some of those bits are flipped.

Viterbi decoding finds the maximum-likelihood transmitted sequence by the most-likely path through the 64-state trellis. For hard decisions the branch metric is the **Hamming distance** between the received pair and each branch's expected $(c1, c2)$; the decoder keeps, per state, the survivor path with the smallest cumulative metric:

$$PM_t(\sigma') = \min \text{ over predecessors } \sigma \quad [PM_{t-1}(\sigma) + \text{Hamming}(r_t, c(\sigma \rightarrow \sigma'))]$$

then traces the survivors back to output the bits. It corrects any error pattern up to roughly $\text{floor}((d_{\text{free}}-1)/2) - d_{\text{free}} = 10$ for this code — per window; validated $\sim 100\%$ CRC-OK up to $\sim 2\%$ raw BER. (Two original bugs — encoder emitting from the *next* state, traceback storing the bit instead of the predecessor state — were fixed.)

Soft-decision Viterbi (Euclidean instead of Hamming metric, using the raw symbol distances) would add ~ 2 dB but isn't implemented.

Common alternatives: LDPC and Turbo codes (near-Shannon, used in Wi-Fi/5G), Reed-Solomon (burst errors), Polar codes (5G control), soft-decision Viterbi (~ 2 dB better than the hard-decision decoder here).

8. Error detection + ARQ

- **CRC-16-CCITT** appended to every chunk; the RX drops any chunk that fails the check (guarantees the delivered message is error-free).
- **Stop-and-wait ARQ** (`--role source_arq / sink_arq`): the source sends a chunk and waits for an ACK; on `--timeout 3000` ms it retransmits; it advances when ACKed and exits when all chunks are ACKed. The sink CRC-verifies each chunk, ACKs only verified ones, re-ACKs duplicates, and auto-stops once every chunk is in.
- **ACK transport** (`--ack-transport`):
 - `tcp` (default) — ACK over a socket; **no reverse RF path needed**. Source connects to `--ack-host 127.0.0.1` : `--ack-port 5599` (the sink listens there). Ideal when both radios share one host.
 - `rf` — ACK sent back over the air on the second RF path (needs a clean reverse link).

Math/theory. Treat the message bits as coefficients of a polynomial $M(x)$ over $GF(2)$; append 16 zeros ($M(x) \cdot x^{16}$) and divide by the CRC-16-CCITT generator $G(x) = x^{16} + x^{12} + x^5 + 1 \pmod{2}$ / XOR long division). The 16-bit remainder is the CRC; the RX recomputes it and flags any nonzero mismatch. Because an undetected error requires the error polynomial $E(x)$ to be an exact multiple of $G(x)$, this structure guarantees detection of: all single- and double-bit errors, any odd number of errors (G has the factor $x+1$), and all burst errors of length ≤ 16 . Residual undetected probability $\approx 2^{-16}$ for random errors.

Stop-and-wait is the simplest ARQ (throughput limited by idling one round-trip per chunk). *Alternatives:* Go-Back-N / Selective-Repeat (pipelined, higher throughput), Hybrid-ARQ (FEC + ARQ combined — what you get here with `--fec` on: FEC corrects most errors, CRC+retransmit catches the rest).

One-way mode (`--role tx / rx`, no ACK): the TX cycles all chunks `--tx-reps` times; the RX reassembles from CRC-verified chunks and auto-stops `--rx-idle-timeout 8` s after the last burst.

9. Automatic gain control (AGC)

`--AGC_type` `Feed` (feed-forward, default) or `Closed` (closed-loop). Feed-forward measures the detected burst's RMS and scales it to unit power in one shot (no loop lag); closed-loop adapts a gain over time. The energy detector's captured burst feeds the AGC before sync.

10. Visualization / EVM

`--viz true` (default) captures one TX and one RX burst and **auto-saves a figure on exit** to `<viz-dir>/<scheme>/figure.png` (per-modulation folder, `--viz-dir viz`). The 2×3 figure shows TX/RX **time domain**, **spectrum** (FFT), and **constellation** with the ideal points overlaid and an **EVM %** readout. Rough EVM ceilings for reliable decode: QPSK tolerates ~30–35% (more with FEC), 16-QAM needs <~12%, 64-QAM <~6%. Render manually with `python3 tools/plot_viz.py <dir> --fs 1.6e6 --save out.png`.

11. Complete command-line reference

Mode / roles

Option	Default	Controls
<code>--role</code>	both	tx, rx, both, source_arq, sink_arq
<code>--mode</code>	source	legacy source/sink selector
<code>--tx-reps</code>	20	one-way: times to cycle all chunks (no ACK)
<code>--rx-idle-timeout</code>	8.0	one-way RX: auto-stop after N s of no bursts (0 = until Ctrl-C)
<code>--num_bits</code>	1000	payload bits per packet
<code>--interval</code>	3000	TX gap between packets (ms)

ARQ / ACK

Option	Default	Controls
<code>--ack-transport</code>	tcp	ACK channel: tcp or rf
<code>--ack-host</code>	127.0.0.1	TCP: sink IP the source connects to
<code>--ack-port</code>	5599	TCP: ACK socket port
<code>--timeout</code>	3000	source ACK timeout (ms) before retransmit
<code>--timer_interval</code>	1000	sink ACK timer interval (ms)

Modulation / waveform / FEC

Option	Default	Controls
<code>--scheme</code>	QPSK	constellation (§2)
<code>--fec</code>	false	rate-1/2 K=7 conv + Viterbi (must match both ends)
<code>--waveform</code>	sc	sc (single-carrier) or ofdm
<code>--ofdm-fft</code>	64	OFDM subcarrier count (FFT size)
<code>--ofdm-cp</code>	16	OFDM cyclic-prefix length
<code>--ofdm-tx-peak</code>	0.5	OFDM peak scaling (PAPR / clipping guard)
<code>--sps</code>	2	samples/symbol (informational)

Preamble / pulse shaping

Option	Default	Controls
<code>--preamble</code>	m-sequence	m-sequence or Zadoff
<code>--m</code>	5	m-sequence order (length $2^m - 1$)
<code>--add_preamble</code>	true	prepend preamble to each packet

Option	Default	Controls
--filter_type	rrc	rrc / rc / lp
--roll_off	0.25	RRC/RC roll-off β
--num_taps	151	filter tap count
--U / --D	2 / 1	pulse-shaper up/down sampling (wire sps = U/D)
--symbol_rate	0.8e6	symbol rate (Hz)
--num_threads	1	FFT threads

Energy detection

Option	Default	Controls
--alpha	0.95	IIR power-smoothing (larger = smoother)
--det-adaptive (--IIR_threshold_adaptive)	true	use noise_floor $\times$ mult vs fixed threshold
--det-mult (--IIR_threshold_multiplier)	5.0	adaptive threshold multiplier (raise OTA to 10–30)
--det-threshold (--energy_threshold)	1e-7	fixed threshold (adaptive off only)
--det-continuous	true	re-track noise floor during idle vs one-shot calib
--energy_packet_size	3300	samples to collect after detection (auto-sized)
--IIR_window_size	20	IIR window size

Synchronization / timing / phase

Option	Default	Controls
--sync-threshold (--sync_threshold)	15.0	ACQ correlation gate (set below true peak, above noise)
--sps_sync	5	samples/symbol at matched-filter output
--recv_msg_len	auto	data symbols to extract per burst (auto from scheme)
--timing_loop_bw	0.015	Gardner TED loop bandwidth BnT
--timing_damping	0.707	Gardner TED damping
--phase_loop_bw	0.02	carrier phase PLL bandwidth
--phase_damping	0.707	carrier phase PLL damping

Equalizer

Option	Default	Controls
--eq_type	None	None / LMS / RLS / DFE (see §6 — LMS diverges OTA)
--eq_taps	11	equalizer tap count
--eq_mu	0.3	NLMS step (DD / real-preamble training)
--eq_dd	false	decision-directed tracking after LS training

RF hardware (TX)

Option	Default	Controls
--tx-args	""	UHD TX device (e.g. serial=30CD424)
--tx-rate	1.6e6	TX sample rate (Hz) = symbol_rate $\cdot$ U/D
--tx-freq	2.412e9	TX centre frequency (Hz)
--tx-gain	20	TX gain (dB)
--tx-bw	1.0e6	TX analog bandwidth (Hz)
--tx-ant	TX/RX	TX antenna port

Option	Default	Controls
--tx-subdev	A:A	TX subdev spec
--tx-channel	0	TX channel index

RF hardware (RX)

Option	Default	Controls
--rx-args	""	UHD RX device (e.g. serial=30CD3F7)
--rx-rate	1.6e6	RX sample rate (Hz), must equal TX rate
--rx-freq	2.412e9	RX centre frequency (Hz)
--rx-gain	30	RX gain (dB) — key for EVM/clipping (§ recipes)
--rx-bw	1.0e6	RX analog bandwidth (Hz)
--rx-ant	RX2	RX antenna port
--rx-subdev	A:A	RX subdev spec
--rx-channel	0	RX channel index

Common RF / buffering / misc

Option	Default	Controls
--ref	internal	clock reference: internal / external / mimo
--settling	0.2	RF settling time (s)
--uhd_timeout	1000	UHD TX timeout (ms)
--samps_per_buff	10000	samples per UHD receive buffer
--num_recv_request	0	total samples to receive (0 = continuous)
--AGC_type	Feed	Feed (feed-forward) or Closed (closed-loop)
--skip-rate-check	off	bypass startup rate-consistency check
--viz	true	capture + auto-save TX/RX figure on exit
--viz-dir	viz	base dir for figures (per-scheme subfolder made)

12. Verified command recipes (915 MHz, VERT900 antennas)

QPSK OFDM + ARQ (reliable OTA demo). Start the sink first.

```
# Sink / RX
./sdr_system --role sink_arq --rx-args serial=30CD3F7 --tx-args serial=30CD3F7 \
  --rx-subdev A:A --rx-ant RX2 --rx-freq 915e6 --rx-rate 1.6e6 --rx-gain 25 \
  --waveform ofdm --ofdm-fft 64 --ofdm-cp 16 --scheme QPSK --fec true \
  --ack-transport tcp --ack-port 5599 --det-mult 3 --viz-dir viz
# Source / TX
./sdr_system --role source_arq --tx-args serial=30CD424 --rx-args serial=30CD424 \
  --tx-subdev A:A --tx-ant TX/RX --tx-freq 915e6 --tx-rate 1.6e6 --tx-gain 82 \
  --waveform ofdm --ofdm-fft 64 --ofdm-cp 16 --ofdm-tx-peak 0.5 --scheme QPSK --fec true \
  --ack-transport tcp --ack-host 127.0.0.1 --ack-port 5599 --timeout 3000 --viz-dir viz
```

Single-carrier QPSK one-way (swap --waveform ofdm ... for --waveform sc, roles rx/tx, add --tx-reps 20).

16-QAM / 64-QAM: same as above with --scheme 16QAM (or 64QAM) on both ends — but needs much higher SNR (EVM < ~12% / 6%). Use a cabled link (SMA + 30–40 dB attenuator) or very close antennas; tune --rx-gain so the RX Peak= stays under ~0.8 (below 1.0 = ADC clipping) while RMS is as high as possible. Best OTA point found so far: --rx-gain 21 --tx-gain 80 --ofdm-tx-peak 0.45.

Level-tuning cheatsheet (the #1 cause of "nothing decodes")

- RX log Peak > 1.0 -> **clipping**, lower --rx-gain (fatal for QAM).
- No detections (bursts=0) -> signal too weak, raise --tx-gain or --rx-gain, or lower --det-mult.
- Detections but CRC fails -> check the constellation EVM in the figure; if the order is too high for the link SNR, drop to QPSK or improve the link.
- Keep --fec true on marginal links; it closes the last few % of bit errors.